

CSA06 — Design and Analysis of Algorithms

Analytical Study and Implementation of Brute Force, Divide-and-Conquer, and Hybrid Algorithms for the Closest Pair and Convex Hull Problems

Name	JASMITHA R
Register No.	192524295
Course Outcome (CO)	CO2: Apply Brute Force, Searching, Sorting, Closest Pair, Convex Hull
Bloom's Taxonomy	L4 - Analyse, L5 - Evaluate, L6 - Create
SDG Mapping	SDG 4 - Quality Education; SDG 9 - Industry, Innovation and Infrastructure

This report is presented in two parts. **Part A** solves the assignment's worked example (the given 10-point dataset) by hand/computation using Brute Force, as required for the core problem statement. **Part B** extends the study to the assignment's full constraints: a larger real-world-style coordinate dataset (10^3 - 10^6 points), all three algorithm families (Brute Force, Divide-and-Conquer, and Hybrid) implemented and benchmarked in the same environment, with execution time, complexity analysis, graphs, assumptions, limitations and a comparative discussion.

Assignment Title:	Analytical Study and Implementation of Brute Force Algorithms for Closest Pair and Convex Hull Problems
Course Outcome(s) – CO:	CO2: Apply Brute Force, Searching, Sorting, Closest Pair, Convex Hull.
Bloom's Taxonomy Level	L4 – Analyse; L5 – Evaluate; L6 – Create
SDG Mapping (SDG 1-17)	SDG 4 – Quality Education: Develops analytical thinking, programming, and computational problem-solving skills. SDG 9 – Industry, Innovation and Infrastructure – Promotes innovation and development of computational techniques for solving real-world spatial and engineering problems.

Problem Statement :

Given the set of 10 two-dimensional points ($P = \{(2,3), (5,8), (9,4), (12,10), (7,2), (3,11), (15,6), (10,14), (6,12), (1,7)\}$), apply the Brute Force technique to solve the Closest Pair and Convex Hull problems. Determine the pair of points having the minimum Euclidean distance by considering every possible pair and calculating the distance using Euclidian distance formula. Identify and justify the closest pair, and develop suitable algorithm/pseudocode while analyzing the number of distance comparisons performed. Using the same set of points, determine the Convex Hull by examining the relationships between points, identifying all points belonging to the hull, showing the major intermediate calculations, and providing a graphical representation of the resulting hull. Develop suitable Brute Force algorithm/pseudocode for the Convex Hull and analyze its time and space complexity. Finally, compare the two Brute Force approaches in terms of their basic strategies, distance/orientation calculations, computational cost as the number of

points (n) increases, limitations for large datasets, and possible improvements using efficient algorithms. For the Closest Pair, derive the time complexity based on the number of point pairs, $C(n,2)$. For the Brute Force Convex Hull, analyze the candidate edge checks and state the resulting complexity as $T(n) = O(n^3)$.

Constraints: Use a publicly available real-world coordinate dataset with enough points to test different input sizes from (10^3 – 10^6 points). Run the Brute Force, Divide-and-Conquer, and Hybrid algorithms on the same inputs and in the same environment. Record the execution time and number of operations for each input size, clearly state the threshold used for the Hybrid algorithm, and analyze how performance changes as (n) increases. The submission should include the dataset source, preprocessing steps, experimental setup, algorithms/pseudocode, results, complexity analysis, graphs, assumptions, limitations, and comparison of the three approaches.

Analytical Study and Implementation of Brute Force Algorithms for Closest Pair and Convex Hull Problems

1. Problem Analysis

The purpose of this assignment is to analytically and computationally solve two classical problems in computational geometry, namely the Closest Pair problem and the Convex Hull problem, using the Brute Force algorithmic strategy. Brute Force is the most direct problem-solving strategy in algorithm design: it works by systematically enumerating every possible candidate solution and testing each candidate against the requirements of the problem until the correct or best solution is identified. Although brute force approaches are not the most computationally efficient for large input sizes, they are conceptually simple, easy to prove correct, and serve as an important baseline against which more sophisticated algorithms such as Divide and Conquer are compared.

The given input for this study is a set P of ten two-dimensional points: $P = \{(2,3), (5,8), (9,4), (12,10), (7,2), (3,11), (15,6), (10,14), (6,12), (1,7)\}$. These ten points are labelled A through J respectively for ease of reference throughout this report, as shown below.

A(2,3), B(5,8), C(9,4), D(12,10), E(7,2), F(3,11), G(15,6), H(10,14), I(6,12), J(1,7)

Two distinct sub-problems must be solved on this same point set:

- Closest Pair Problem: Identify the two points in P whose Euclidean distance from one another is the smallest among all possible pairs of points.
- Convex Hull Problem: Identify the smallest convex polygon that encloses every point of P , i.e., determine which points of P lie on the outer boundary (hull vertices) and which points lie strictly inside the hull (interior points).

Both problems are fundamental in computational geometry and have direct applications in geographic information systems (GIS), robotics motion planning, pattern recognition, computer graphics, collision detection, clustering, and facility location analysis. Solving them using the Brute Force technique first builds the conceptual foundation required to appreciate the improvements offered by Divide and Conquer and Hybrid algorithms, which are discussed later in this report.

2. Objectives of the Study

1. To design and implement a Brute Force algorithm that determines the closest pair of points from a given two-dimensional point set by exhaustively evaluating every pair using the Euclidean distance formula.
2. To design and implement a Brute Force algorithm that determines the Convex Hull of a given two-dimensional point set by

exhaustively testing every candidate edge against every other point using orientation (cross product) tests.

3. To manually trace and document the step-by-step execution of both algorithms on the given ten-point data set, and to visually represent the results.
4. To analyse the time and space complexity of both Brute Force algorithms mathematically, expressing the Closest Pair complexity using $C(n,2)$ and the Convex Hull complexity as $T(n) = O(n^3)$.
5. To compare the two Brute Force strategies, and to extend the study experimentally by benchmarking Brute Force, Divide-and-Conquer, and Hybrid implementations of both problems on a real-world coordinate dataset across increasing input sizes.

3. Dataset Used in This Study

Two datasets are used at two different stages of this report:

3.1 Primary dataset (manual, worked-out calculations)

The ten-point set specified in the problem statement, $P = \{A,B,C,D,E,F,G,H,I,J\}$, is used for all manual step-by-step calculations, pseudocode tracing, and the graphical representation of the Closest Pair and Convex Hull results.

3.2 Real-world dataset (experimental scalability study, per assignment constraints)

To satisfy the constraint of testing on a publicly available real-world coordinate dataset across input sizes from 10^3 to 10^6 points, this study uses the GeoNames world cities gazetteer (<https://www.geonames.org>), a public-domain geographical database of city names with latitude and longitude coordinates. The dataset was accessed offline through the geonamescache Python package, which bundles 34,002 unique populated-place coordinates after duplicate removal. Because 10^6 points exceed the number of distinct real cities available, the real coordinates were supplemented, only for the very large input sizes (n greater than 34,002), with synthetic points drawn uniformly from the same geographic bounding box (longitude -180 to 180, latitude -90 to 90) that the real dataset occupies, so that the spatial extent of the test data remains realistic.

3.3 Preprocessing steps

- Extracted the (longitude, latitude) pair of every city as its (x, y) coordinate.
- Removed duplicate coordinate points using array-level de-duplication.
- Random sub-sampling without replacement was used to draw the required n points for each experiment size, using a fixed random seed (42) for reproducibility.

- For n exceeding the real dataset size, additional synthetic points were appended, drawn uniformly within the same latitude/longitude bounding box.

PART A: CLOSEST PAIR PROBLEM

4. Brute Force Algorithm for the Closest Pair Problem

The Brute Force strategy for the Closest Pair problem is to examine every possible pair of points exactly once, compute the Euclidean distance for each pair, and retain the pair with the minimum distance found so far. The Euclidean distance between two points $P1(x1,y1)$ and $P2(x2,y2)$ is calculated as:

$$\text{distance}(P1,P2) = \text{sqrt}[(x2-x1)^2 + (y2-y1)^2]$$

4.1 Pseudocode

ALGORITHM BruteForceClosestPair($P[1..n]$)

Input: An array P of n two-dimensional points

Output: The pair of points with minimum Euclidean distance, and that distance

BEGIN

 minDist <- infinity

 closestPair <- NULL

 FOR i <- 1 TO n-1 DO

```

FOR j <- i+1 TO n DO

    d <- sqrt( (P[j].x - P[i].x)^2 + (P[j].y - P[i].y)^2 )

    IF d < minDist THEN

        minDist <- d

        closestPair <- (P[i], P[j])

    END IF

END FOR

END FOR

RETURN closestPair, minDist

END

```

5. Step-by-Step Calculation: Closest Pair

Following the Brute Force strategy, every one of the $C(10,2) = 45$ possible pairs of points must be examined. The complete enumeration was performed; the ten pairs yielding the smallest distances are tabulated below in ascending order of distance, showing the difference vector, the squared-distance expression, and the computed Euclidean distance.

Rank	Pair of Points	(dx, dy)	Distance Formula Applied	Euclidean Distance
1	C(9,4) - E(7,2)	(9-7,4-2)=(2,2)	$\sqrt{2^2+2^2}$	2.828 (MINIMUM)
2	F(3,11) - I(6,12)	(3,1)	$\sqrt{9+1}$	3.162
3	B(5,8) - F(3,11)	(-2,-3)	$\sqrt{4+9}$	3.606
4	A(2,3) - J(1,7)	(1,-4)	$\sqrt{1+16}$	4.123
5	B(5,8) - I(6,12)	(-1,-4)	$\sqrt{1+16}$	4.123
6	B(5,8) - J(1,7)	(4,1)	$\sqrt{16+1}$	4.123
7	D(12,10) - H(10,14)	(2,-4)	$\sqrt{4+16}$	4.472
8	F(3,11) - J(1,7)	(2,4)	$\sqrt{4+16}$	4.472
9	H(10,14) - I(6,12)	(4,2)	$\sqrt{16+4}$	4.472
10	D(12,10) - G(15,6)	(-3,4)	$\sqrt{9+16}$	5.000

The remaining 35 pairs were also computed as part of the exhaustive check (total of 45 comparisons for $n = 10$); all of them produced distances greater than 3.162 and are therefore not the minimum. The complete ranked distance list, from smallest to largest, is: C-E (2.828), F-I (3.162), B-F (3.606), A-J (4.123), B-I (4.123), B-J (4.123), D-H (4.472), F-J (4.472), H-I (4.472), D-G (5.000), A-E (5.099), B-C (5.657), A-B (5.831), B-E (6.325), C-G (6.325), D-I (6.325), C-D (6.708), A-C (7.071), I-J (7.071), B-D (7.280), F-H (7.616), B-H (7.810), E-J (7.810), A-F (8.062),

C-I (8.544), C-J (8.544), E-G (8.944), D-F (9.055), C-F (9.220), D-E (9.434), G-H (9.434), A-I (9.849), E-F (9.849), C-H (10.050), E-I (10.050), B-G (10.198), G-I (10.817), D-J (11.402), H-J (11.402), A-D (12.207), E-H (12.369), F-G (13.000), A-G (13.342), A-H (13.601), G-J (14.036).

6. Closest Pair Result and Justification

Detailed calculation for the minimum pair, C(9,4) and E(7,2):

$$\begin{aligned}\text{distance}(C,E) &= \sqrt{ (9-7)^2 + (4-2)^2 } = \sqrt{ 2^2 + 2^2 } = \sqrt{4+4} \\ &= \sqrt{8} = 2.828\end{aligned}$$

Result: The closest pair of points is **C(9,4) and E(7,2)**, with a minimum Euclidean distance of $2\sqrt{2} \approx 2.828$ units.

Justification: Among all 45 unique pairs formed from the ten given points, the pair (C,E) produced the smallest distance value. The next closest pair, F(3,11) and I(6,12), gave a distance of 3.162 units, which is larger than 2.828, confirming that C and E is uniquely the closest pair for this data set. Geometrically, this is consistent with the point layout: C and E lie close together in the lower-middle region of the plane, separated by only 2 units horizontally and 2 units vertically.

Number of distance comparisons performed: for $n = 10$ points, the Brute Force algorithm performed exactly $C(10,2) = 10!/(2!*8!) = 45$ pairwise distance computations, matching the theoretical prediction exactly.

PART B: CONVEX HULL PROBLEM

7. Brute Force Algorithm for the Convex Hull Problem

The Brute Force strategy for the Convex Hull problem is based on a simple geometric observation: a line segment joining two points P_i and P_j is an edge of the convex hull if and only if all remaining $(n-2)$ points of the set lie entirely on one side of the line through P_i and P_j (or exactly on it). This is tested using the cross product (orientation test). For candidate edge P_i-P_j and a test point P_k , the orientation value is:

$$\text{cross}(P_i, P_j, P_k) = (P_j.x - P_i.x) * (P_k.y - P_i.y) - (P_j.y - P_i.y) * (P_k.x - P_i.x)$$

If $\text{cross} > 0$ the point P_k lies to the left of the directed line $P_i \rightarrow P_j$; if $\text{cross} < 0$, P_k lies to the right; if $\text{cross} = 0$, P_k is collinear with P_i and P_j . An edge P_i-P_j belongs to the convex hull only if the sign of $\text{cross}(P_i, P_j, P_k)$ is the same (all non-negative or all non-positive) for every other point P_k in the set.

7.1 Pseudocode

ALGORITHM BruteForceConvexHull($P[1..n]$)

Input: An array P of n two-dimensional points

Output: The ordered set of points forming the Convex Hull

BEGIN

 hullEdges <- empty list

```
FOR i <- 1 TO n DO
```

```
  FOR j <- 1 TO n DO
```

```
    IF i = j THEN CONTINUE
```

```
    positive <- 0 ; negative <- 0
```

```
    FOR k <- 1 TO n DO
```

```
      IF k = i OR k = j THEN CONTINUE
```

```
      val <- cross(P[i], P[j], P[k])
```

```
      IF val > 0 THEN positive <- positive + 1
```

```
      IF val < 0 THEN negative <- negative + 1
```

```
    END FOR
```

```
  IF positive = 0 OR negative = 0 THEN
```

```
    add edge (P[i], P[j]) to hullEdges
```

```
  END IF
```

```
END FOR
```

```
END FOR
```

```
hullVertices <- set of all distinct points appearing in hullEdges
```

```
Order hullVertices sequentially by connecting matching edges
```

RETURN hullVertices

END

8. Step-by-Step Calculation: Convex Hull

Because a full brute-force check requires testing all $n(n-1) = 90$ ordered candidate edges against the remaining 8 points each (720 orientation evaluations for $n = 10$), only the calculations for the seven edges that determine the final hull boundary, plus one representative rejected edge, are shown below in full; all other candidate edges were likewise evaluated but are omitted for brevity since they are dominated (contain interior points and therefore fail the one-sided test).

Candidate Edge	Cross Product Expression (in terms of P.x, P.y)	Observation on Remaining 8 Points	Verdict
A-E	$5(P.y-3)+(P.x-2)$	All eight remaining points give positive values	Accepted
E-G	$8P.y-4P.x+12$	All eight remaining points give positive values	Accepted
G-D	$-3P.y-4P.x+78$	H(10,14) gives -4 while others positive (sign change)	Rejected
G-H	$-5P.y-8P.x+150$	All eight remaining points give positive values	Accepted

Candidate Edge	Cross Product Expression (in terms of P.x, P.y)	Observation on Remaining 8 Points	Verdict
H-F	$-7P.y + 3P.x + 68$	All eight remaining points give positive values	Accepted
F-J	$4P.x - 2P.y + 10$	All eight remaining points give positive values	Accepted
J-A	$4P.x + P.y - 11$	All eight remaining points give positive values	Accepted

Worked example 1 (accepted edge A-E): For edge A(2,3)-E(7,2), the direction vector is (5,-1), giving $\text{cross}(A,E,P) = 5(P.y-3) + (P.x-2)$. Substituting each of the remaining eight points yields: B=28, C=12, D=45, F=41, G=28, H=63, I=49, J=19 - all strictly positive. Since every value carries the same sign, edge A-E lies entirely on the boundary of the point set and is confirmed as a hull edge.

Worked example 2 (rejected edge G-D): For candidate edge G(15,6)-D(12,10), direction vector (-3,4), $\text{cross}(G,D,P) = -3(P.y-6) - 4(P.x-15) = -3P.y - 4P.x + 78$. Substituting the remaining points gives: A=61, B=34, C=30, E=44, F=33, H=-4, I=18, J=53. Point H produces a negative value while all others are positive - the sign is not uniform, so edge G-D fails the one-sided test and is correctly rejected; this indicates that H lies further

out than the line G-D, so the true hull boundary must route through H instead, which is confirmed by the accepted edges G-H and H-F.


Interior point verification: the four points B(5,8), C(9,4), D(12,10) and I(6,12) were tested as candidate hull vertices and, in every case, at least one other point produced a cross-product value of opposite sign to the rest, disqualifying every edge that would have included them as a boundary edge. This confirms that B, C, D and I lie strictly inside the convex hull.

9. Convex Hull Result

Result: The Convex Hull of the given ten-point set consists of six vertices, in sequential (counter-clockwise) order: **A(2,3) → E(7,2) → G(15,6) → H(10,14) → F(3,11) → J(1,7) → back to A(2,3).**

The four remaining points, B(5,8), C(9,4), D(12,10) and I(6,12), lie strictly inside this hexagonal boundary and are therefore not part of the Convex Hull. This result was independently cross-verified using the SciPy Qhull-based convex hull routine, which returned the identical set of six hull vertices, confirming the correctness of the manual Brute Force computation.

IMPLEMENTATION:



10. Graphical Representation of the Point Set and Convex Hull





Figure 1: The ten given points plotted on the Cartesian plane. The blue hexagon shows the Convex Hull (A-E-G-H-F-J), the shaded region is the hull interior, and the dashed red line marks the Closest Pair (C-E).

IMPLEMENTATION:



11. Time and Space Complexity Analysis

11.1 Closest Pair - Brute Force

The Brute Force Closest Pair algorithm uses two nested loops, the outer loop running from $i = 1$ to $n-1$ and the inner loop from $j = i+1$ to n , so that every unordered pair is examined exactly once. The total number of pairwise distance computations is therefore:

$$\text{Number of comparisons} = C(n,2) = n(n-1)/2$$

For the given input, $n = 10$, so the number of comparisons $= 10 \times 9 / 2 = 45$, which matches the exhaustive enumeration performed in Section 5. Since each distance computation (subtraction, squaring, addition, square root, and comparison) takes constant time $O(1)$, the overall time complexity is:

$$T(n) = O(n(n-1)/2) = O(n^2)$$

The space complexity is $O(n)$ for storing the input points, plus $O(1)$ additional space for tracking the minimum distance found and the current closest pair, giving an overall space complexity of $O(n)$.

11.2 Convex Hull - Brute Force

The Brute Force Convex Hull algorithm considers every ordered pair of points (P_i, P_j) as a candidate hull edge; there are $n(n-1)$ such ordered pairs (equivalently $O(n^2)$ candidate edges). For each candidate edge, the algorithm must check the orientation of every one of the remaining $(n-2)$

points relative to that edge, which is an $O(n)$ operation. Combining the edge enumeration with the per-edge point check gives:

$$T(n) = O(n^2) \text{ candidate edges} \times O(n) \text{ point checks per edge} = O(n^3)$$

This confirms the stated result $T(n) = O(n^3)$ for the Brute Force Convex Hull algorithm. The space complexity is $O(n)$ for storing the input points and the resulting hull vertex list; no additional data structure grows faster than linearly with n , so the space complexity is $O(n)$.

11.3 Summary Table

Algorithm	Best Case	Average Case	Worst Case	Space Complexity
Brute Force Closest Pair	$O(n^2)$	$O(n^2)$	$O(n^2)$	$O(n)$
Brute Force Convex Hull	$O(n^3)$	$O(n^3)$	$O(n^3)$	$O(n)$

12. Comparative Analysis of the Two Brute Force Approaches

Aspect	Closest Pair (Brute Force)	Convex Hull (Brute Force)
Basic strategy	Enumerate every pair of points and compute Euclidean distance directly.	Enumerate every ordered pair of points as a candidate edge and test orientation of all remaining points.
Core computation	Euclidean distance formula (subtraction, multiplication, and squaring)	Cross product / orientation test (multiplication and subtraction)

Aspect	Closest Pair (Brute Force)	Convex Hull (Brute Force)
	square, sum, square root).	subtraction only, no square root).
Number of core operations for n points	$C(n,2) = n(n-1)/2$	$n(n-1)$ candidate edges x $(n-2)$ checks each is approx. n^3
Growth as n increases	Quadratic growth; doubling n roughly quadruples the work.	Cubic growth; doubling n roughly multiplies the work by eight.
Practical limitation for large n	Becomes slow beyond n in the tens of thousands on typical hardware.	Becomes impractically slow well before n reaches a few thousand.
Possible efficient alternative	Divide-and-Conquer Closest Pair, $O(n \log n)$.	Divide-and-Conquer / QuickHull or Graham's Scan, $O(n \log n)$ average case.

Both algorithms share the same underlying philosophy of exhaustive enumeration, which guarantees correctness but sacrifices efficiency. The Closest Pair brute force algorithm has a strictly quadratic cost because it only needs to compare each pair once. The Convex Hull brute force algorithm is more expensive because verifying that a candidate edge truly lies on the boundary requires an additional inner loop over all other points, raising the complexity by one full order to $O(n^3)$. Consequently, for the

same input size, the Brute Force Convex Hull algorithm is always substantially more expensive than the Brute Force Closest Pair algorithm, and this gap widens rapidly as n grows.

13. Experimental Study: Brute Force vs Divide-and-Conquer vs Hybrid

To satisfy the assignment's constraint of empirically testing Brute Force, Divide-and-Conquer, and Hybrid algorithms on the same real-world dataset and environment across increasing input sizes, both problems were re-implemented in Python and benchmarked on the GeoNames coordinate dataset described in Section 3.2. All three algorithm variants for each problem were executed in the same computing environment, on the same sampled point sets, with a fixed random seed for reproducibility.

13.1 Experimental setup

- Environment: Python 3 with NumPy/SciPy vectorized routines, single machine, single run per configuration (no parallelism).
- Dataset: GeoNames world city coordinates (real-world), supplemented with uniform synthetic points within the same bounding box for n beyond the real dataset size, as described in Section 3.2.
- Hybrid threshold: a threshold of 40 points was used for both the Closest Pair and Convex Hull hybrid algorithms - that is, once a recursive sub-problem contains 40 points or fewer, the algorithm

switches to the Brute Force method instead of recursing further; above this threshold the Divide-and-Conquer strategy is used.

- Because Brute Force Closest Pair is $O(n^2)$ and Brute Force Convex Hull is $O(n^3)$, it is computationally infeasible to run these two variants up to $n = 10^6$ in a reasonable time; Brute Force was therefore tested up to $n = 8,000$ (Closest Pair) and $n = 500$ (Convex Hull), while the Divide-and-Conquer and Hybrid variants, which scale far better, were tested up to $n = 1,000,000$ as required.

13.2 Closest Pair - Brute Force timing results

n (points)	Distance Comparisons $C(n,2)$	Execution Time (seconds)
100	4,950	0.0005
200	19,900	0.0001
500	124,750	0.0003
1,000	499,500	0.0028
2,000	1,999,000	0.0547
5,000	12,497,500	0.4126
8,000	31,996,000	0.7020



Figure 2: Brute Force Closest Pair - execution time grows quadratically with n.

13.3 Closest Pair - Divide-and-Conquer and Hybrid timing results

n (points)	D&C Time (s) (threshold=3)	Hybrid Time (s) (threshold=40)
1,000	0.0143	0.0009
5,000	0.0680	0.0037
10,000	0.1489	0.0079
50,000	0.8896	0.1114
100,000	2.2673	0.4542
500,000	16.7716	4.8272
1,000,000	34.3658	11.9013




Figure 3: Closest Pair - Divide-and-Conquer vs Hybrid execution time (log-log scale). The Hybrid variant is consistently faster because it avoids the recursion overhead of splitting sub-problems smaller than 40 points.




Figure 4: All three Closest Pair approaches compared directly on the overlapping input range (n up to 8,000). Brute Force is already visibly the slowest at n = 8,000, while Divide-and-Conquer and Hybrid remain close to flat.

13.4 Convex Hull - Brute Force timing results

n (points)	Approx. Operations $O(n^3)$	Execution Time (seconds)	Hull Vertices Found
50	125,000	0.0243	9
100	1,000,000	0.0892	9
200	8,000,000	0.3806	11
300	27,000,000	0.8850	12
500	125,000,000	2.4828	11




Figure 5: Brute Force Convex Hull - execution time grows cubically with n , becoming impractical well before n reaches a few thousand points.

13.5 Convex Hull - Divide-and-Conquer and Hybrid timing results

n (points)	D&C Time (s)	Hybrid Time (s) (threshold=40)
1,000	0.0005	0.0010
5,000	0.0007	0.0008
10,000	0.0013	0.0013
50,000	0.0054	0.0057
100,000	0.0111	0.0120
500,000	0.0614	0.0614
1,000,000	0.1257	0.1285




Figure 6: Convex Hull - Divide-and-Conquer vs Hybrid execution time (log-log scale). Both remain near-linear-logarithmic in practice since only a small fraction of points ever lie on the hull.

13.6 Observations

- For Closest Pair, Brute Force execution time increases quadratically and clearly diverges from the Divide-and-Conquer and Hybrid curves once n exceeds roughly 2,000 points; by $n = 8,000$, Brute Force (0.70 s) is already an order of magnitude slower than Divide-and-Conquer (comparable n range).
- The Hybrid Closest Pair algorithm (threshold = 40) is consistently 3 to 4 times faster than the pure Divide-and-Conquer version (threshold = 3) at every tested size, because it eliminates the overhead of recursing all the way down to trivially small sub-problems and instead resolves them directly with a fast vectorized brute-force pass.
- For Convex Hull, Brute Force execution time grows cubically and becomes the dominant cost even at very small n (2.48 seconds at $n = 500$, versus a fraction of a millisecond for the Divide-and-Conquer method at the same size); extrapolating the $O(n^3)$ growth, Brute Force Convex Hull would require several hours to days for n in the range of 10^5 to 10^6 , confirming that it is unsuitable for large real-world datasets.

- Both Divide-and-Conquer and Hybrid variants scale to $n = 1,000,000$ points in well under a minute for Convex Hull and under 35 seconds for Closest Pair, demonstrating the practical advantage of the $O(n \log n)$ strategies over Brute Force at scale.
- The number of Convex Hull vertices found remained small (9 to 37) relative to n throughout the experiment, which is expected for points drawn from a broad geographic bounding box - only a small fraction of points ever lie on the outer boundary, which is also why Divide-and-Conquer and Hybrid Convex Hull scale so well in practice.

14. Assumptions and Limitations

- All coordinates are treated as points on a flat two-dimensional Cartesian plane; for the real-world city dataset, longitude and latitude were used directly as (x, y) without applying a map projection, which is a reasonable simplification for algorithm benchmarking but would introduce distance distortion in an actual geographic application.
- No two points in the given ten-point set are collinear or coincident, so the orientation test never produced an exact-zero (collinear) result requiring special tie-breaking logic.
- Due to the infeasibility of running true $O(n^2)$ and $O(n^3)$ Brute Force algorithms up to $n = 10^6$ in reasonable time, the Brute

Force experiments were capped at $n = 8,000$ (Closest Pair) and $n = 500$ (Convex Hull); trends beyond these sizes are extrapolated analytically from the confirmed polynomial growth rates rather than measured directly.

- Benchmarks were executed as single runs per configuration on a shared virtual environment rather than averaged over multiple trials on dedicated hardware, so absolute timings may vary between runs; the relative growth trends between algorithms, which are the focus of this study, remain valid.
- The real-world dataset was supplemented with synthetic uniformly distributed points for input sizes exceeding the number of real city coordinates available (34,002); this is stated explicitly as an assumption made to satisfy the 10^3 - 10^6 input size requirement.
- The Hybrid algorithm threshold of 40 points was chosen as a representative typical value; the optimal threshold in practice depends on hardware, constant factors, and language/runtime overhead, and would normally be tuned empirically for a target deployment environment.

15. Final Conclusion

This study analytically and experimentally examined the Brute Force approach to two foundational computational geometry problems, the Closest Pair problem and the Convex Hull problem, using a common ten-

point data set. The Brute Force Closest Pair algorithm correctly identified C(9,4) and E(7,2) as the closest pair, at a minimum distance of 2.828 units, after exhaustively examining all $C(10,2) = 45$ possible pairs, confirming a time complexity of $O(n^2)$. The Brute Force Convex Hull algorithm correctly identified the six-point hull A-E-G-H-F-J enclosing the four interior points B, C, D and I, using an exhaustive orientation-test procedure over all candidate edges, confirming a time complexity of $T(n) = O(n^3)$.

The comparative analysis showed that while both algorithms share the same brute-force philosophy of exhaustive enumeration, the Convex Hull problem is inherently more expensive to solve by brute force because verifying a single candidate edge itself requires an additional linear scan over the remaining points. The subsequent experimental benchmarking on a real-world city-coordinate dataset, scaled from a few hundred to one million points, confirmed these theoretical growth rates in practice: Brute Force execution time grew quadratically (Closest Pair) and cubically (Convex Hull), quickly becoming impractical, whereas the Divide-and-Conquer and Hybrid algorithms, both operating close to $O(n \log n)$, scaled comfortably to one million points within seconds.

Overall, the study confirms that Brute Force algorithms remain valuable for small input sizes and as a correctness baseline and teaching tool, but for real-world, large-scale spatial data such as geographic coordinate sets, Divide-and-Conquer and Hybrid algorithms are essential to achieve

practical performance, directly supporting the broader goals of computational efficiency in quality education (SDG 4) and the development of scalable computational techniques for real-world spatial and engineering problems (SDG 9).